

Hadoop Cluster Assignment

AKSHAYA NAGARAJAN

SJSU ID: 014535014

STEPS FOR CONFIGURING HADOOP CLUSTER:

Install brew on mac – using the command :

/usr/bin/ruby -e "\$(curl -fsSL

<https://raw.githubusercontent.com/Homebrew/install/master/install>

```

Last login: Sat Sep 28 16:44:45 from ::1
AkshayaNagarajanMacbookPro:~ akshayanagarajan$ /usr/bin/ruby -e "$(curl -fsSL https://raw.githubusercontent.com/Homebrew/install/master/install)"
--> This script will install:
/usr/local/bin/brew
/usr/local/share/doc/homebrew
/usr/local/share/man/man1/brew.1
/usr/local/share/zsh/site-functions/_brew
/usr/local/etc/bash_completion.d/brew
/usr/local/Homebrew

Press RETURN to continue or any other key to abort

```

Installed Homebrew

```

AkshayaNagarajanMacbookPro:~ akshayanagarajan$ /usr/bin/ruby -e "$(curl -fsSL https://raw.githubusercontent.com/Homebrew/install/master/install)"
--> This script will install:
/usr/local/bin/brew
/usr/local/share/doc/homebrew
/usr/local/share/man/man1/brew.1
/usr/local/share/zsh/site-functions/_brew
/usr/local/etc/bash_completion.d/brew
/usr/local/Homebrew

Press RETURN to continue or any other key to abort
--> Downloading and installing Homebrew...
HEAD is now at db18713c8 Merge pull request #6496 from Homebrew/dependabot/bundler/docs/rake-13.0.0
Already up-to-date.
--> Installation successful!

--> Homebrew has enabled anonymous aggregate formulae and cask analytics.
Read the analytics documentation (and how to opt-out) here:
  https://docs.brew.sh/Analytics

--> Homebrew is run entirely by unpaid volunteers. Please consider donating:
  https://github.com/Homebrew/brew#donations

--> Next steps:
- Run `brew help` to get started
- Further documentation:
  https://docs.brew.sh

```

Checking the Homebrew version

```

AkshayaNagarajanMacbookPro:bin akshayanagarajan$ brew --version
Homebrew 2.1.11
Homebrew/homebrew-core (git revision 0e28; last commit 2019-09-28)
AkshayaNagarajanMacbookPro:bin akshayanagarajan$

```

Installing Hadoop using homebrew

```
AkshayaNagarajanMacbookPro:~ akshayanagarajan$ brew install hadoop
Warning: hadoop 3.1.2 is already installed and up-to-date
To reinstall 3.1.2, run `brew reinstall hadoop`
AkshayaNagarajanMacbookPro:~ akshayanagarajan$ brew reinstall hadoop
--> Reinstalling hadoop
--> Downloading https://www.apache.org/dyn/closer.cgi?path=hadoop/common/hadoop-3.1.2/hadoop-3.1.2.tar.gz
Already downloaded: /Users/akshayanagarajan/Library/Caches/Homebrew/downloads/190eebdb181ba6d46acd05da5df80ee0ff76d392324f9805ca7f702093c2834--hadoop-3.1.2.tar.gz
🔗 /usr/local/Cellar/hadoop/3.1.2: 21,686 files, 774.1MB, built in 32 seconds
Removing: /usr/local/Cellar/hadoop/hdfs... (63 files, 5MB)
```

Hadoop version

```
AkshayaNagarajanMacbookPro:~ akshayanagarajan$ hadoop version
Hadoop 3.1.2
Source code repository https://github.com/apache/hadoop.git -r 1019dde65bcf12e05ef48ac71e84550d589e5d9a
Compiled by sunilg on 2019-01-29T01:39Z
Compiled with protoc 2.5.0
From source with checksum 64b8bdd4ca6e77cce75a93eb09ab2a9
This command was run using /usr/local/Cellar/hadoop/3.1.2/libexec/share/hadoop/common/hadoop-common-3.1.2.jar
AkshayaNagarajanMacbookPro:~ akshayanagarajan$ |
```

Configuring Hadoop-env.sh in /usr/local/Cellar/hadoop/3.1.2/libexec/etc/hadoop

```
###
# HDFS Balancer specific parameters
###
# Specify the JVM options to be used when starting the HDFS Balancer.
# These options will be appended to the options specified as HADOOP_OPTS
# and therefore may override any similar flags set in HADOOP_OPTS
#
# export HDFS_BALANCER_OPTS=""

###
# HDFS Mover specific parameters
###
# Specify the JVM options to be used when starting the HDFS Mover.
# These options will be appended to the options specified as HADOOP_OPTS
# and therefore may override any similar flags set in HADOOP_OPTS
#
# export HDFS_MOVER_OPTS=""

###
# Router-based HDFS Federation specific parameters
# Specify the JVM options to be used when starting the RBF Routers.
# These options will be appended to the options specified as HADOOP_OPTS
# and therefore may override any similar flags set in HADOOP_OPTS
#
# export HDFS_DFSROUTER_OPTS=""
###

###
# Advanced Users Only!
###

#
# When building Hadoop, one can add the class paths to the commands
# via this special env var:
# export HADOOP_ENABLE_BUILD_PATHS="true"

#
# To prevent accidents, shell commands be (superficially) locked
# to only allow certain users to execute certain subcommands.
# It uses the format of (command)_(subcommand)_USER.
#
# For example, to limit who can execute the namenode command,
# export HDFS_NAMENODE_USER=hdfs

export HADOOP_OPTS="$HADOOP_OPTS -Djava.net.preferIPv4Stack=true -Djava.security.krb5.realm= -Djava.security.krb5.kdc="
export JAVA_HOME="/Library/Java/JavaVirtualMachines/jdk1.8.0_221.jdk/Contents/Home"
```

Configuring Core-site.xml in /usr/local/Cellar/hadoop/3.1.2/libexec/etc/hadoop

```
<?xml version="1.0" encoding="UTF-8"?>
<?xml-stylesheet type="text/xsl" href="configuration.xsl"?>
<!--
  Licensed under the Apache License, Version 2.0 (the "License");
  you may not use this file except in compliance with the License.
  You may obtain a copy of the License at

    http://www.apache.org/licenses/LICENSE-2.0

  Unless required by applicable law or agreed to in writing, software
  distributed under the License is distributed on an "AS IS" BASIS,
  WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied.
  See the License for the specific language governing permissions and
  limitations under the License. See accompanying LICENSE file.
-->

<!-- Put site-specific property overrides in this file. -->

<configuration>
  <property>
    <name>hadoop.tmp.dir</name>
    <value>/usr/local/Cellar/hadoop/hdfs/tmp</value>
    <description>A base for other temporary directories.</description>
  </property>
  <property>
    <name>fs.default.name</name>
    <value>hdfs://localhost:8020</value>
  </property>
</configuration>
```

Configuring Mapred-site.xml in /usr/local/Cellar/hadoop/3.1.2/libexec/etc/hadoop

```
<?xml version="1.0"?>
<?xml-stylesheet type="text/xsl" href="configuration.xsl"?>
<!--
  Licensed under the Apache License, Version 2.0 (the "License");
  you may not use this file except in compliance with the License.
  You may obtain a copy of the License at

    http://www.apache.org/licenses/LICENSE-2.0

  Unless required by applicable law or agreed to in writing, software
  distributed under the License is distributed on an "AS IS" BASIS,
  WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied.
  See the License for the specific language governing permissions and
  limitations under the License. See accompanying LICENSE file.
-->

<!-- Put site-specific property overrides in this file. -->

<configuration>
  <property>
    <name>mapred.job.tracker</name>
    <value>localhost:8021</value>
  </property>
</configuration>
~
~
~
~
~
```

Configuring Hdfs-site.xml in

/usr/local/Cellar/hadoop/3.1.2/libexec/etc/hadoop

```
<?xml version="1.0" encoding="UTF-8"?>
<?xml-stylesheet type="text/xsl" href="configuration.xsl"?>
<!--
  Licensed under the Apache License, Version 2.0 (the "License");
  you may not use this file except in compliance with the License.
  You may obtain a copy of the License at

    http://www.apache.org/licenses/LICENSE-2.0

  Unless required by applicable law or agreed to in writing, software
  distributed under the License is distributed on an "AS IS" BASIS,
  WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied.
  See the License for the specific language governing permissions and
  limitations under the License. See accompanying LICENSE file.
-->

<!-- Put site-specific property overrides in this file. -->

<configuration>
  <property>
    <name>dfs.replication</name>
    <value>1</value>
  </property>
</configuration>
~
~
~
~
~
```

Configuring Yarn-site.xml in

/usr/local/Cellar/hadoop/3.1.2/libexec/etc/hadoop

```
<?xml version="1.0"?>
<!--
  Licensed under the Apache License, Version 2.0 (the "License");
  you may not use this file except in compliance with the License.
  You may obtain a copy of the License at

    http://www.apache.org/licenses/LICENSE-2.0

  Unless required by applicable law or agreed to in writing, software
  distributed under the License is distributed on an "AS IS" BASIS,
  WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied.
  See the License for the specific language governing permissions and
  limitations under the License. See accompanying LICENSE file.
-->
<configuration>
<!-- Site specific YARN configuration properties -->
  <property>
    <name>yarn.nodemanager.aux-services</name>
    <value>mapreduce_shuffle</value>
  </property>
  <property>
    <name>yarn.nodemanager.env-whitelist</name>
    <value>JAVA_HOME,HADOOP_COMMON_HOME,HADOOP_HDFS_HOME,HADOOP_CONF_DIR,CLASSPATH_PREPEND_DISTCACHE,HADOOP_YARN_HOME,HADOOP_MAPRED_HOME</value>
  </property>
</configuration>
```

Generate SSH key

```

AkshayaNagarajanMacbookPro:~ akshayanagarajan$ ssh-keygen -t rsa
Generating public/private rsa key pair.
Enter file in which to save the key (/Users/akshayanagarajan/.ssh/id_rsa):
Created directory '/Users/akshayanagarajan/.ssh'.
Enter passphrase (empty for no passphrase):
Enter same passphrase again:
Your identification has been saved in /Users/akshayanagarajan/.ssh/id_rsa.
Your public key has been saved in /Users/akshayanagarajan/.ssh/id_rsa.pub.
The key fingerprint is:
SHA256:82dfDX7imWqIt2c0dW0XbUTJNbWocRD7fnAU5NnU6X0 akshayanagarajan@AkshayaNagarajanMacbookPro
The key's randomart image is:
+---[RSA 2048]---+
|      o. o*&|
|      o o*X|
|      o o.*=|
|      = o.E|
|      S . +.+o|
|      o +.o..|
|      ..oooo.+|
|      . oo+o.* |
|      ..=..= |
+-----[SHA256]-----+

```

Copying the generated SSH key to authorized_keys

```

AkshayaNagarajanMacbookPro:~ akshayanagarajan$ cat ~/.ssh/id_rsa.pub >> ~/.ssh/authorized_keys

```

Checking if the SSH localhost works

```

AkshayaNagarajanMacbookPro:hadoop akshayanagarajan$ ssh localhost
Last login: Sat Sep 28 19:46:06 2019

```


Formatting HDFS using the command: hdfs namenode –format

```
Default (set)  #1  Default (set)  #2
WARNING: /usr/local/Cellar/hadoop/3.1.2/libexec/logs does not exist. Creating.
2019-09-28 16:51:51,486 INFO namenode.NameNode: STARTUP MSG:
/*****
STARTUP MSG: Starting NameNode
STARTUP MSG: host = akshayanagarajanmacbookpro/192.168.1.149
STARTUP MSG: args = [-format]
STARTUP MSG: version = 3.1.2
STARTUP MSG: classpath = /usr/local/Cellar/hadoop/3.1.2/libexec/etc/hadoop:/usr/local/Cellar/hadoop/3.1.2/libexec/share/hadoop/common/lib/jaxb-impl-2.2.3-1.jar:/usr/local/Cellar/hadoop/3.1.2/libexec/share/hadoop/common/lib/hadoop-auth-3.1.2.jar:/usr/local/Cellar/hadoop/3.1.2/libexec/share/hadoop/common/lib/slf4j-log4j12-1.7.25.jar:/usr/local/Cellar/hadoop/3.1.2/libexec/share/hadoop/common/lib/jetty-server-9.3.24.v20180605.jar:/usr/local/Cellar/hadoop/3.1.2/libexec/share/hadoop/common/lib/jetty-util-9.3.24.v20180605.jar:/usr/local/Cellar/hadoop/3.1.2/libexec/share/hadoop/common/lib/woodstox-core-5.0.3.jar:/usr/local/Cellar/hadoop/3.1.2/libexec/share/hadoop/common/lib/htrace-core4-4.1.0-incubating.jar:/usr/local/Cellar/hadoop/3.1.2/libexec/share/hadoop/common/lib/kerby-xdm-1.0.1.jar:/usr/local/Cellar/hadoop/3.1.2/libexec/share/hadoop/common/lib/snappy-java-1.0.5.jar:/usr/local/Cellar/hadoop/3.1.2/libexec/share/hadoop/common/lib/jetty-io-9.3.24.v20180605.jar:/usr/local/Cellar/hadoop/3.1.2/libexec/share/hadoop/common/lib/hadoop-annotations-3.1.2.jar:/usr/local/Cellar/hadoop/3.1.2/libexec/share/hadoop/common/lib/slf4j-api-1.7.25.jar:/usr/local/Cellar/hadoop/3.1.2/libexec/share/hadoop/common/lib/jackson-jaxrs-1.9.13.jar:/usr/local/Cellar/hadoop/3.1.2/libexec/share/hadoop/common/lib/commons-logging-1.1.3.jar:/usr/local/Cellar/hadoop/3.1.2/libexec/share/hadoop/common/lib/kerby-pkix-1.0.1.jar:/usr/local/Cellar/hadoop/3.1.2/libexec/share/hadoop/common/lib/avro-1.7.7.jar:/usr/local/Cellar/hadoop/3.1.2/libexec/share/hadoop/common/lib/jetty-webapp-9.3.24.v20180605.jar:/usr/local/Cellar/hadoop/3.1.2/libexec/share/hadoop/common/lib/jackson-annotation-2.7.8.jar:/usr/local/Cellar/hadoop/3.1.2/libexec/share/hadoop/common/lib/javax-servlet-api-3.1.0.jar:/usr/local/Cellar/hadoop/3.1.2/libexec/share/hadoop/common/lib/kerb-server-1.0.1.jar:/usr/local/Cellar/hadoop/3.1.2/libexec/share/hadoop/common/lib/kerby-krb5-1.0.1.jar:/usr/local/Cellar/hadoop/3.1.2/libexec/share/hadoop/common/lib/kerby-config-1.0.1.jar:/usr/local/Cellar/hadoop/3.1.2/libexec/share/hadoop/common/lib/jcip-annotations-1.0.1.jar:/usr/local/Cellar/hadoop/3.1.2/libexec/share/hadoop/common/lib/kerby-asn1-1.0.1.jar:/usr/local/Cellar/hadoop/3.1.2/libexec/share/hadoop/common/lib/netty-3.10.5.Final.jar:/usr/local/Cellar/hadoop/3.1.2/libexec/share/hadoop/common/lib/jetty-ee11-api-1.1.jar:/usr/local/Cellar/hadoop/3.1.2/libexec/share/hadoop/common/lib/jackson-xc-1.9.13.jar:/usr/local/Cellar/hadoop/3.1.2/libexec/share/hadoop/common/lib/zookeeper-3.4.13.jar:/usr/local/Cellar/hadoop/3.1.2/libexec/share/hadoop/common/lib/kerb-identity-1.0.1.jar:/usr/local/Cellar/hadoop/3.1.2/libexec/share/hadoop/common/lib/hittclint-4.4.5.2.jar:/usr/local/Cellar/hadoop/3.1.2/libexec/share/hadoop/common/lib/kerby-config-1.0.1.jar:/usr/local/Cellar/hadoop/3.1.2/libexec/share/hadoop/common/lib/guava-11.0.2.jar:/usr/local/Cellar/hadoop/3.1.2/libexec/share/hadoop/common/lib/json-smart-2.3.jar:/usr/local/Cellar/hadoop/3.1.2/libexec/share/hadoop/common/lib/kerb-util-1.0.1.jar:/usr/local/Cellar/hadoop/3.1.2/libexec/share/hadoop/common/lib/curator-framework-2.13.0.jar:/usr/local/Cellar/hadoop/3.1.2/libexec/share/hadoop/common/lib/curator-client-2.13.0.jar:/usr/local/Cellar/hadoop/3.1.2/libexec/share/hadoop/common/lib/jersey-servlet-1.19.jar:/usr/local/Cellar/hadoop/3.1.2/libexec/share/hadoop/common/lib/jetty-security-9.3.24.v20180605.jar:/usr/local/Cellar/hadoop/3.1.2/libexec/share/hadoop/common/lib/jaxb-api-2.2.11.jar:/usr/local/Cellar/hadoop/3.1.2/libexec/share/hadoop/common/lib/commons-io-2.5.jar:/usr/local/Cellar/hadoop/3.1.2/libexec/share/hadoop/common/lib/commons-configuration2-2.11.jar:/usr/local/Cellar/hadoop/3.1.2/libexec/share/hadoop/common/lib/jackson-core-2.7.8.jar:/usr/local/Cellar/hadoop/3.1.2/libexec/share/hadoop/common/lib/curator-recipes-2.13.0.jar:/usr/local/Cellar/hadoop/3.1.2/libexec/share/hadoop/common/lib/kerb-crypto-1.0.1.jar:/usr/local/Cellar/hadoop/3.1.2/libexec/share/hadoop/common/lib/jsp-api-2.1.jar:/usr/local/Cellar/hadoop/3.1.2/libexec/share/hadoop/common/lib/kerb-crypto-1.0.1.jar:/usr/local/Cellar/hadoop/3.1.2/libexec/share/hadoop/common/lib/commons-codec-1.11.jar:/usr/local/Cellar/hadoop/3.1.2/libexec/share/hadoop/common/lib/asn-5.0.4.jar:/usr/local/Cellar/hadoop/3.1.2/libexec/share/hadoop/common/lib/kerb-admin-1.0.1.jar:/usr/local/Cellar/hadoop/3.1.2/libexec/share/hadoop/common/lib/audience-annotations-0.5.0.jar:/usr/local/Cellar/hadoop/3.1.2/libexec/share/hadoop/common/lib/paranamer-2.3.jar:/usr/local/Cellar/hadoop/3.1.2/libexec/share/hadoop/common/lib/kerb-client-1.0.1.jar:/usr/local/Cellar/hadoop/3.1.2/libexec/share/hadoop/common/lib/commons-collections-3.2.2.jar:/usr/local/Cellar/hadoop/3.1.2/libexec/share/hadoop/common/lib/jsp-api-2.1.jar:/usr/local/Cellar/hadoop/3.1.2/libexec/share/hadoop/common/lib/jetty-http-9.3.24.v20180605.jar:/usr/local/Cellar/hadoop/3.1.2/libexec/share/hadoop/common/lib/curator-recipes-2.13.0.jar:/usr/local/Cellar/hadoop/3.1.2/libexec/share/hadoop/common/lib/commons-compress-1.18.jar:/usr/local/Cellar/hadoop/3.1.2/libexec/share/hadoop/common/lib/token-provider-1.0.1.jar:/usr/local/Cellar/hadoop/3.1.2/libexec/share/hadoop/common/lib/jetty-servlet-9.3.24.v20180605.jar:/usr/local/Cellar/hadoop/3.1.2/libexec/share/hadoop/common/lib/jersey-core-1.19.jar:/usr/local/Cellar/hadoop/3.1.2/libexec/share/hadoop/common/lib/jsch-0.1.54.jar:/usr/local/Cellar/hadoop/3.1.2/libexec/share/hadoop/common/lib/kerb-common-1.0.1.jar:/usr/local/Cellar/hadoop/3.1.2/libexec/share/hadoop/common/lib/httpcore-4.4.4.jar:/usr/local/Cellar/hadoop/3.1.2/libexec/share/hadoop/common/lib/kerb-simplekdc-1.0.1.jar:/usr/local/Cellar/hadoop/3.1.2/libexec/share/hadoop/common/lib/jersey-server-1.19.jar:/usr/local/Cellar/hadoop/3.1.2/libexec/share/hadoop/common/lib/accessors-smart-1.2.jar:/usr/local/Cellar/hadoop/3.1.2/libexec/share/hadoop/common/lib/net3-1.1.jar:/usr/local/Cellar/hadoop/3.1.2/libexec/share/hadoop/common/lib/gson-2.4.jar:/usr/local/Cellar/hadoop/3.1.2/libexec/share/hadoop/common/lib/commons-lang-2.6.jar:/usr/local/Cellar/hadoop/3.1.2/libexec/share/hadoop/common/lib/kerb-core-1.0.1.jar:/usr/local/Cellar/hadoop/3.1.2/libexec/share/hadoop/common/lib/commons-net-3.6.jar:/usr/local/Cellar/hadoop/3.1.2/libexec/share/hadoop/common/lib/kerby-util-1.0.1.jar:/usr/local/Cellar/hadoop/3.1.2/libexec/share/hadoop/common/lib/jackson-databind-2.7.8.jar:/usr/local/Cellar/hadoop/3.1.2/libexec/share/hadoop/common/lib/jackson-core-2.7.8.jar:/usr/local/Cellar/hadoop/3.1.2/libexec/share/hadoop/common/lib/jetty-servlet-9.3.24.v20180605.jar:/usr/local/Cellar/hadoop/3.1.2/libexec/share/hadoop/common/lib/commons-beanutils-1.9.3.jar:/usr/local/Cellar/hadoop/3.1.2/libexec/share/hadoop/common/lib/commons-lang3-3.4.jar:/usr/local/Cellar/hadoop/3.1.2/libexec/share/hadoop/common/lib/jsr305-3.0.0.jar:/usr/local/Cellar/hadoop/3.1.2/libexec/share/hadoop/common/lib/nimbus-jose-jwt-4.4.1.jar:/usr/local/Cellar/hadoop/3.1.2/libexec/share
2019-09-28 16:51:52,340 INFO namenode.FSDirectory: POSIX ACL inheritance enabled? true
2019-09-28 16:51:52,340 INFO namenode.FSDirectory: XAttrs enabled? true
2019-09-28 16:51:52,340 INFO namenode.NameNode: Caching file names occurring more than 10 times
2019-09-28 16:51:52,350 INFO snapshot.SnapshotManager: Loaded config captureOpenFiles: false, skipCaptureAccessTimeOnlyChange: false, snapshotDisallowSnapRootDescendant: true, maxSnapshotsLimit: 65536
2019-09-28 16:51:52,360 INFO snapshot.SnapshotManager: SkipList is disabled
2019-09-28 16:51:52,364 INFO util.GSet: Computing capacity for map cachedBlocks
2019-09-28 16:51:52,364 INFO util.GSet: VM type = 64-bit
2019-09-28 16:51:52,364 INFO util.GSet: 0.25% max memory 4 GB = 10.2 MB
2019-09-28 16:51:52,364 INFO util.GSet: capacity = 2^20 = 1048576 entries
2019-09-28 16:51:52,384 INFO metrics.TopMetrics: NNTop conf: dfs.namenode.top.window.num.buckets = 10
2019-09-28 16:51:52,384 INFO metrics.TopMetrics: NNTop conf: dfs.namenode.top.num.entries = 10
2019-09-28 16:51:52,384 INFO metrics.TopMetrics: NNTop conf: dfs.namenode.top.windows.minutes = 1,5,25
2019-09-28 16:51:52,389 INFO namenode.FSNamesystem: Retry cache on namenode is enabled
2019-09-28 16:51:52,389 INFO namenode.FSNamesystem: Retry cache will use 0.03 of total heap and retry cache entry expiry time is 6000000 millis
2019-09-28 16:51:52,391 INFO util.GSet: Computing capacity for map NameNodeRetryCache
2019-09-28 16:51:52,391 INFO util.GSet: VM type = 64-bit
2019-09-28 16:51:52,392 INFO util.GSet: 0.02999999329447746% max memory 4 GB = 1.2 MB
2019-09-28 16:51:52,392 INFO util.GSet: capacity = 2^17 = 131072 entries
2019-09-28 16:51:52,424 INFO namenode.FSImage: Allocated new BlockPoolId: BP-368296384-192.168.1.149-156074712417
2019-09-28 16:51:52,456 INFO namenode.Storage: Storage directory /usr/local/Cellar/hadoop/hdfs/tmp/dfs/name has been successfully formatted.
2019-09-28 16:51:52,467 INFO namenode.FSImageFormatProtobuf: Saving image file /usr/local/Cellar/hadoop/hdfs/tmp/dfs/name/current/fsimage.ckpt_00000000000000000000 using no compression
2019-09-28 16:51:52,552 INFO namenode.FSImageFormatProtobuf: Image file /usr/local/Cellar/hadoop/hdfs/tmp/dfs/name/current/fsimage.ckpt_00000000000000000000 of size 403 bytes saved in 0 seconds
2019-09-28 16:51:52,577 INFO namenode.NMStorageRetentionManager: Going to retain 1 images with txid >= 0
2019-09-28 16:51:52,587 INFO namenode.NameNode: SHUTDOWN_MSG:
/*****
SHUTDOWN_MSG: Shutting down NameNode at akshayanagarajanmacbookpro/192.168.1.149
*****/
AkshayaNagarajanMacbookPro:bin akshayanagarajan$ cd ..
-bash: cd.: command not found
AkshayaNagarajanMacbookPro:bin akshayanagarajan$ ls
container-executor  hadoop  mapred  test-container-executor yarn
AkshayaNagarajanMacbookPro:bin akshayanagarajan$ cd ..
-bash: cd.: command not found
AkshayaNagarajanMacbookPro:bin akshayanagarajan$ cd sbin
AkshayaNagarajanMacbookPro:bin akshayanagarajan$ ls
bin  hadoop-config.sh  hdfs-config.sh  logs  sbin  yarn-config.sh
etc  hadoop-functions.sh  libexec  mapred-config.sh  share
AkshayaNagarajanMacbookPro:libexec akshayanagarajan$ cd sbin
AkshayaNagarajanMacbookPro:sbin akshayanagarajan$ ls
FederationStateStore  hadoop-daemons.sh  mr-jobhistory-daemon.sh  start-balancer.sh  start-yarn.sh  stop-dfs.sh  workers.sh
distribute-exclude.sh  hdfs.sh  refresh-namenodes.sh  start-dfs.sh  stop-all.sh  stop-secure-dns.sh  yarn-daemon.sh
hadoop-daemon.sh  kms.sh  start-all.sh  start-secure-dns.sh  stop-balancer.sh  stop-yarn.sh  yarn-daemons.sh
AkshayaNagarajanMacbookPro:sbin akshayanagarajan$ ./start-all.sh
WARNING: Attempting to start all Apache Hadoop daemons as akshayanagarajan in 10 seconds.
WARNING: This is not a recommended production deployment configuration.
WARNING: Use CTRL-C to abort.
```

Start hadoop using the command : ./start-all.sh after changing the directory to /usr/local/Cellar/hadoop/3.1.2/sbin/

```
AkshayaNagarajanMacbookPro:sbin akshayanagarajan$ ./start-all.sh
WARNING: Attempting to start all Apache Hadoop daemons as akshayanagarajan in 10 seconds.
WARNING: This is not a recommended production deployment configuration.
WARNING: Use CTRL-C to abort.
Starting namenodes on [localhost]
Starting datanodes
Starting secondary namenodes [AkshayaNagarajanMacbookPro]
2019-09-28 19:19:56,673 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Starting resourcemanager
Starting nodemanagers
```

```
AkshayaNagarajanMacbookPro:sbin akshayanagarajan$ ./start-all.sh
WARNING: Attempting to start all Apache Hadoop daemons as akshayanagarajan in 10 seconds.
WARNING: This is not a recommended production deployment configuration.
WARNING: Use CTRL-C to abort.
Starting namenodes on [localhost]
Starting datanodes
Starting secondary namenodes [AkshayaNagarajanMacbookPro]
2019-09-28 19:19:56,673 WARN util.NativeCodeLoader: Unable to load native-hadoop library
Starting resourcemanager
Starting nodemanagers
```

Check the hadoop daemons running on the machine

```
AkshayaNagarajanMacbookPro:sbin akshayanagarajan$ jps
38528 DataNode
38962 NodeManager
39030 Jps
38424 NameNode
38667 SecondaryNameNode
38862 ResourceManager
```

Overview 'localhost:8020' (active)

Started:	Sat Sep 28 19:19:49 -0700 2019
Version:	3.1.2, r1019d9e6bdc12e05ef48ac71e84550d589e5d9a
Compiled:	Mon Jan 28 17:39:00 -0800 2019 by sunlig from branch-3.1.2
Cluster ID:	CID-99B2b8f-3356-4d79-8042-34529f05e25
Block Pool ID:	BP-1980556455-192.168.1.149-1569719815844

Summary

Security is off.
Safemode is off.
1 files and directories, 0 blocks (0 replicated blocks, 0 erasure coded block groups) = 1 total filesystem object(s).
Heap Memory used 107.32 MB of 446.5 MB Heap Memory. Max Heap Memory is 3.56 GB.
Non Heap Memory used 46.83 MB of 48.71 MB Committed Non Heap Memory. Max Non Heap Memory is -unbounded-.

Configured Capacity:	233.47 GB
Configured Remote Capacity:	0 B
DFS Used:	12 KB (0%)
Non DFS Used:	63.39 GB
DFS Remaining:	167.44 GB (71.72%)
Block Pool Used:	12 KB (0%)
DataNodes usages% (Min/Median/Max/stdDev):	0.00% / 0.00% / 0.00% / 0.00%
Live Nodes	1 (Decommissioned: 0, In Maintenance: 0)

Summary

Security is off.
Safemode is off.
1 files and directories, 0 blocks (0 replicated blocks, 0 erasure coded block groups) = 1 total filesystem object(s).
Heap Memory used 107.32 MB of 446.5 MB Heap Memory. Max Heap Memory is 3.56 GB.
Non Heap Memory used 46.83 MB of 48.71 MB Committed Non Heap Memory. Max Non Heap Memory is -unbounded-.

Configured Capacity:	233.47 GB
Configured Remote Capacity:	0 B
DFS Used:	12 KB (0%)
Non DFS Used:	63.39 GB
DFS Remaining:	167.44 GB (71.72%)
Block Pool Used:	12 KB (0%)
DataNodes usages% (Min/Median/Max/stdDev):	0.00% / 0.00% / 0.00% / 0.00%
Live Nodes	1 (Decommissioned: 0, In Maintenance: 0)
Dead Nodes	0 (Decommissioned: 0, In Maintenance: 0)
Decommissioning Nodes	0
Entering Maintenance Nodes	0
Total Datanode Volume Failures	0 (0 B)
Number of Under-Replicated Blocks	0
Number of Blocks Pending Deletion	0
Block Deletion Start Time	Sat Sep 28 19:19:49 -0700 2019
Last Checkpoint Time	Sat Sep 28 19:19:49 -0700 2019

NameNode Journal Status

localhost:9870/dfshealth.html#tab-overview

192.168.1.1

Decommissioning Nodes0

Entering Maintenance Nodes0

Total Datanode Volume Failures0 (0 B)

Number of Under-Replicated Blocks0

Number of Blocks Pending Deletion0

Block Deletion Start TimeSat Sep 28 19:19:49 -0700 2019

Last Checkpoint TimeSat Sep 28 19:19:49 -0700 2019

NameNode Journal Status

Current transaction ID: 25

Journal ManagerState

FileJournalManager(root=/usr/local/Cellar/hadoop/hdfs/tmp/dfs/name) EditLogFileOutputStream(/usr/local/Cellar/hadoop/hdfs/tmp/dfs/name/current/edits_inprogress_00000000000000000025)

NameNode Storage

Storage DirectoryTypeState

/usr/local/Cellar/hadoop/hdfs/tmp/dfs/nameIMAGE_AND_EDITSActive

DFS Storage Types

Storage TypeConfigured CapacityCapacity UsedCapacity RemainingBlock Pool UsedNodes In Service

DISK233.47 GB12 KB (6%)167.44 GB (71.72%)12 KB1

Hadoop, 2018.

localhost:9870/dfshealth.html#tab-overview

192.168.1.1

Decommissioning Nodes0

Entering Maintenance Nodes0

Total Datanode Volume Failures0 (0 B)

Number of Under-Replicated Blocks0

Number of Blocks Pending Deletion0

Block Deletion Start TimeSat Sep 28 19:19:49 -0700 2019

Last Checkpoint TimeSat Sep 28 19:19:49 -0700 2019

NameNode Journal Status

Current transaction ID: 25

Journal ManagerState

FileJournalManager(root=/usr/local/Cellar/hadoop/hdfs/tmp/dfs/name) EditLogFileOutputStream(/usr/local/Cellar/hadoop/hdfs/tmp/dfs/name/current/edits_inprogress_00000000000000000025)

NameNode Storage

Storage DirectoryTypeState

/usr/local/Cellar/hadoop/hdfs/tmp/dfs/nameIMAGE_AND_EDITSActive

DFS Storage Types

Storage TypeConfigured CapacityCapacity UsedCapacity RemainingBlock Pool UsedNodes In Service

DISK233.47 GB12 KB (6%)167.44 GB (71.72%)12 KB1

Hadoop, 2018.

localhost:8086/cluster

192.168.1.1

hadoop

All Applications

Logged in as: dr.who

Cluster

Cluster Metrics

Cluster Nodes Metrics

Scheduler Metrics

Scheduler

Tools

Apps Submitted0

Apps Pending0

Apps Running0

Apps Completed0

Containers Running0

Memory Used0 B

Memory Total8 GB

Memory Reserved0 B

VCoers Used0

VCoers Total8

VCoers Reserved0

Active Nodes1

Decommissioning Nodes0

Decommissioned Nodes0

Lost Nodes0

Unhealthy Nodes0

Rebooted Nodes0

Shutdown Nodes0

Scheduler TypeCapacity Scheduler

Scheduling Resource Type[memory-mb (unit-M), vcores]

Minimum Allocation<memory:1024, vCores:1>

Maximum Allocation<memory:8192, vCores:4>

Maximum Cluster Application Priority0

Show 20 entries

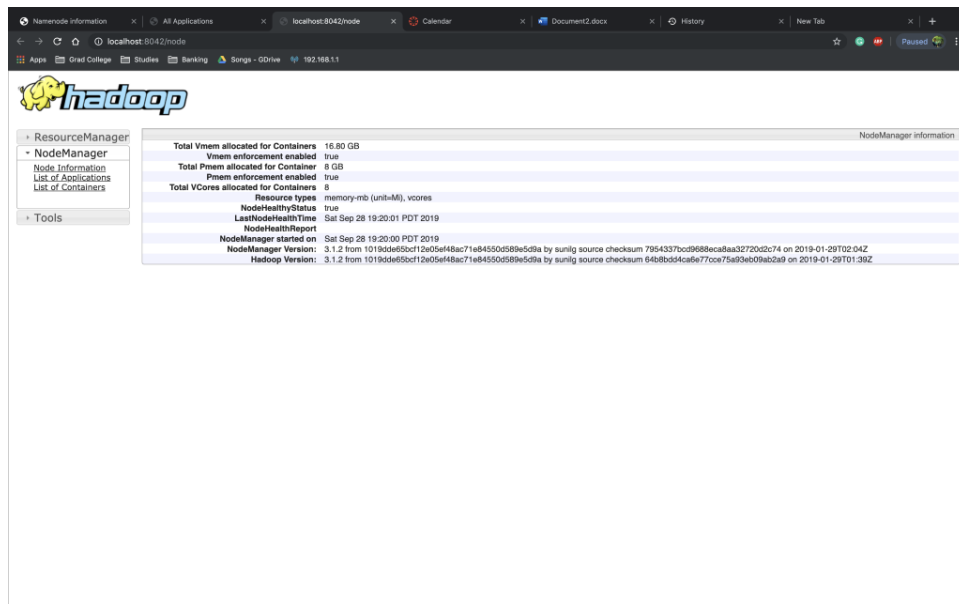
Search

IDUserNameApplicationQueueApplication PriorityStartTimeFinishTimeStateFinalStatusRunningContainersAllocated CPU VCoersAllocated Memory MBReserved CPU VCoersReserved Memory MB% of Queue% of ClusterProgressTracking UIBlacklisted Nodes

No data available in table

Showing 0 to 0 of 0 entries

FirstPreviousNextLast



Stopping Hadoop using the command: `./stop-all.sh` after moving to the directory `/usr/local/Cellar/hadoop/3.1.2/sbin/`

```
AkshayaNagarajanMacbookPro:sbin akshayanagarajan$ ./stop-all.sh
WARNING: Stopping all Apache Hadoop daemons as akshayanagarajan in 10 seconds.
WARNING: Use CTRL-C to abort.
Stopping namenodes on [localhost]
Stopping datanodes
Stopping secondary namenodes [AkshayaNagarajanMacbookPro]
2019-09-28 21:03:05,904 WARN util.NativeCodeLoader: Unable to load native-hadoop li
Stopping nodemanagers
Stopping resourcemanager
```